

Does Confidence-Gating Work?

Empirical Evaluation of a Core AI Oversight Assumption

Rushabh Mistry

Independent Researcher

rushabhmistry16@gmail.com · ORCID: 0009-0001-0568-3993

<https://github.com/Rushhaabhhh/confidence-gating-evaluation>

July 2026

DOI: 10.5281/zenodo.21311853

Abstract

Three 2026 governance frameworks (Singapore’s Model AI Governance Framework for Agentic AI, the U.S. NIST/NCCoE concept paper on agent identity and authorisation, and the EU AI Act, Articles 14–15) require “meaningful human oversight” of autonomous AI agents. A common implementation assumption is that agents can flag low-confidence decisions for human review. We test this assumption empirically. To our knowledge, this is among the first evaluations of confidence-gating measured directly against the oversight assumptions of these frameworks; the underlying observation that language models are poorly calibrated on difficult tasks is consistent with prior work. Verbalized confidence is elicited from two models (Llama 3.3 70B and Gemini 2.5 Flash) across four domains ($N=2,004$ decisions). For the models and benchmarks evaluated, the assumption holds on structured question answering: flagging the least-confident 10% of decisions finds errors at $1.7\text{--}2.9\times$ the base rate. On an open-ended software-engineering judgment task (SWE-bench Verified difficulty estimation), it provides little to no value: lift ≈ 1.0 for both models, and stated confidence occupies so narrow a band ($\sigma < 0.06$) that fixed thresholds are ineffective. Post-hoc recalibration reduces aggregate calibration error by more than 50% in every condition, yet does not restore per-decision discrimination. In a preliminary comparison on SWE-bench ($n=110$), self-consistency (three samples with majority-vote agreement) raises lift to 2.06, suggesting that multi-sample strategies warrant further investigation as oversight mechanisms for open-ended agent tasks. A prototype HMAC-signed audit log demonstrates the tamper-evidence requirement these frameworks also specify. Within this experimental scope, the governance implication is that confidence-gating requirements which do not distinguish by task type risk failing on the open-ended judgments characteristic of many deployed agent settings.

1. Introduction

Autonomous AI agents are being deployed in software engineering, finance, and operations faster than the oversight mechanisms to govern them. Three major 2026 governance instruments independently formalise two requirements: agents must support *meaningful human oversight* and produce *tamper-evident audit trails* (IMDA, 2026; NIST/NCCoE, 2026; EU, 2024). Calls for greater visibility into deployed agents make the same case from the research side (Chan et al., 2024).

A natural and widely assumed implementation of the oversight requirement is **confidence-gating**: the agent flags any decision whose stated confidence falls below a threshold, directing limited human attention to the decisions most likely to be wrong (Galileo, 2026; Parasuraman & Manzey, 2010). This assumption appears in practitioner guides, evaluation pipelines, and proposed technical standards for agent deployment.

What has not been measured, in the context these frameworks care about, is whether the assumption holds. Do current models produce confidence estimates that actually identify their errors? Do those estimates work on the *kinds of tasks* agents perform in practice, such as open-ended, multi-step judgments, or only on structured question answering?

This paper reports such a measurement for two widely deployed models across four benchmarks. We make

three contributions:

- (1) **An empirical test of the confidence-gating assumption** across four task types and two models, with bootstrap confidence intervals and permutation tests. On the three structured QA domains, the assumption largely holds for the models tested. On software-engineering difficulty judgment, it does not.
- (2) **A preliminary comparison with self-consistency** on agent-style tasks. Where verbal confidence yields lift ≈ 1.0 on SWE-bench, vote-based self-consistency over three samples yields lift = 2.06 on the same items ($n=110$).
- (3) **A prototype audit log** that implements the tamper-evidence requirement in the Python standard library using standard cryptographic constructions, with a field-by-field mapping to regulatory clauses.

2. Why This Test Matters

The governance frameworks do not run experiments. The IMDA framework (IMDA, 2026) incorporates feedback from more than 60 companies (version 1.5, May 2026) but provides principles, not measurements. NIST’s approach is “explicitly facilitative rather than prescriptive” (CSA, 2026): the agency identifies gaps and co-invests with industry but has not benchmarked specific oversight mechanisms. The EU AI Act (Art. 14(4)(b)) requires systems to support awareness

Table 1: Proposed audit-record fields and the requirements each satisfies.

Field	Requirement
<code>step_id</code> , <code>timestamp</code>	IMDA: lifecycle monitoring
<code>agent_id</code>	NCCoE: attribute actions to non-human entity
<code>action</code> , <code>input_hash</code>	IMDA: record agent action and intent; hash limits data exposure
<code>confidence</code>	Art. 14(4)(b): automation-bias awareness
<code>oversight_flagged</code>	Art. 14(4)(d): trigger for human override
<code>prev_hash</code>	NCCoE: tamper-evidence via chain integrity
<code>signature</code>	NCCoE: non-repudiation

of automation bias (Parasuraman & Manzey, 2010) and (Art. 14(4)(d)) to enable human override, but specifies neither how to measure confidence nor what threshold to use.

The research community has established that LLMs are systematically overconfident (Guo et al., 2017; Tian et al., 2023; Xiong et al., 2024) and that verbalized confidence is often better calibrated than token probabilities for models tuned with human feedback (Ouyang et al., 2022; Tian et al., 2023; Kadavath et al., 2022). Concurrent work on coding agents finds severe overconfidence across pre-, mid-, and post-execution stages (Kadour et al., 2026). Selective prediction (Geifman & El-Yaniv, 2017), conformal prediction (Vovk et al., 2005), and sampling-based uncertainty estimators (Manakul et al., 2023; Kuhn et al., 2023) offer richer frameworks for assessing when a model can and cannot be trusted, and studies of human-AI teams show that effective reliance depends on humans knowing when to trust the model (Bansal et al., 2021). To our knowledge, however, these methods have not been evaluated against the specific oversight requirements of the 2026 frameworks, or on the task types those frameworks target.

What has not been measured is whether confidence provides a *useful oversight trigger*: not just a calibrated number, but a number that actually finds errors, and whether this varies by task type in ways that matter for how governance requirements should be written.

3. Regulatory Requirements and Audit Record

Table 1 maps each field of the proposed per-decision audit record to the specific regulatory clause it satisfies. This is the paper’s governance engineering contribution: translating three sets of prose requirements into one concrete data structure that a compliance team could build against.

The `confidence` and `oversight_flagged` fields are useful only if the confidence signal is informative. The rest of this paper tests that.

4. Method

4.1 Tasks

We evaluate across four domains spanning qualitatively different task structures. **SWE-bench Verified** (Jimenez et al., 2024) contains genuine GitHub

issues; models estimate $p(\text{fix} > 15 \text{ min})$ against human difficulty labels. This is an *external difficulty judgment*: the model predicts a human annotation, not its own correctness. **GSM8K** (Cobbe et al., 2021) provides math word problems with exact numeric answers. **MMLU** (Hendrycks et al., 2021) samples academic MCQs across 20 subjects. **TruthfulQA** (Lin et al., 2022) provides adversarial factual questions. On the three QA domains, models assess $p(\text{own answer is correct})$, a genuinely introspective quantity. This construct difference matters: the SWE-bench null result may partly reflect that models are poor difficulty estimators, not only poor introspectors.

4.2 Models and elicitation

Llama 3.3 70B (Groq) and Gemini 2.5 Flash (Google AI Studio) are evaluated via public APIs; no fine-tuning is performed. Each item receives a single probability elicitation (0–100) at temperature 0, with a fixed random seed (42) for item sampling. An earlier design eliciting a binary prediction plus a separate confidence rating led to uniformly high confidence regardless of direction, silently disabling the oversight gate; a few-shot prompt variant was also ablated for Llama 3.3 70B (results in the repository).

For self-consistency (Wang et al., 2023), the same items are queried three times at temperature 0.3; the fraction voting for the majority answer serves as the confidence estimate. This is a simplified variant of Self-CheckGPT (Manakul et al., 2023), adapted for binary classification tasks.

4.3 Metrics

Oversight lift is the primary metric: the error rate among the least-confident 10% of decisions divided by the overall base error rate. $\text{Lift} > 1$ means low confidence reliably identifies errors; $\text{lift} \approx 1$ means the gate adds no value over random selection. **AUC-RC** (Geifman & El-Yaniv, 2017) (area under the risk-coverage curve) measures whether sorting by confidence ranks decisions by correctness globally. Both metrics are reported with percentile bootstrap 95% CIs (2,000 resamples, seed 42) and permutation p -values (10,000 permutations) for domain comparisons. The three domain contrasts against SWE-bench are planned comparisons; we report unadjusted p -values and note that all three also survive a Bonferroni correction at $\alpha=0.05$ (adjusted threshold 0.0167).

We also report ECE (Guo et al., 2017) and Brier score (Brier, 1950; Gneiting & Raftery, 2007) before and after three post-hoc recalibration methods (temperature scaling, Platt scaling (Platt, 1999), and isotonic regression (Zadrozny & Elkan, 2002)), each evaluated five-fold out-of-sample. Confidence-in-own-answer standard deviation σ flags *confidence collapse* when $\sigma < 0.10$.

5. Results

5.1 Confidence-gating works on structured QA but fails on SWE-bench

Figure 1 shows the main result. On SWE-bench, both models produce $\text{lift} \approx 1$: the least-confident 10% of decisions are only marginally more error-prone than aver-

age (Llama 1.03, Gemini 1.29). Flagging those items provides essentially no oversight value. On GSM8K, MMLU, and TruthfulQA, lift is substantially above 1 for the most part: 2.86, 2.06, and 2.41 for Llama; 0.82[†], 1.91, and 1.67 for Gemini.

The domain dependence is statistically robust within this sample. The Llama 70B GSM8K-SWE-bench lift difference is +1.83 (95% CI [+1.03, +2.63], permutation $p < 0.001$); MMLU-SWE-bench is +1.02 ([+0.14, +1.90], $p = 0.014$); TruthfulQA-SWE-bench is +1.38 ([+0.32, +2.45], $p = 0.009$). All three contrasts survive Bonferroni correction. A two-proportion z -test finds no significant accuracy difference between models on SWE-bench ($\Delta = -0.008$, $p = 0.90$), which argues against a purely model-specific explanation.

[†] Gemini’s GSM8K lift of 0.82 reflects complete confidence collapse ($\sigma = 0.000$): every item received 100% confidence, so the flagged decile has no different error rate from average. This collapse was observed under the default elicitation prompt and may be sensitive to prompt format; it should be read as an illustrative failure mode of one model-benchmark-prompt combination rather than a general property.

The reliability diagrams (Figure 3) confirm the calibration picture. On GSM8K, Llama achieves 69.2% accuracy but reports 98.6% mean confidence (ECE = 0.294). Gemini reports 100% confidence on every item ($\sigma = 0.000$, ECE = 0.440). On MMLU and TruthfulQA both models are better calibrated.

The confidence distributions (Figure 4) explain why the gate fails: $\sigma < 0.10$ in every condition. The distributions of correct and incorrect decisions overlap almost entirely on SWE-bench. No fixed threshold cleanly separates them; the same cutoff flags 8.7% of one model’s decisions and 0% of another’s, not because the models differ in reliability but because their distributions sit at different absolute levels.

5.2 Post-hoc calibration fixes aggregate error but not the gate

Figure 6 shows that post-hoc recalibration substantially reduces ECE. Platt scaling reduces Llama’s SWE-bench ECE from 0.201 to 0.015 (five-fold out-of-sample). In every condition, at least one method reduces ECE by more than 50% relative. Aggregate calibration is technically recoverable without retraining (Guo et al., 2017; Tian et al., 2023).

However, this does not restore the oversight gate. Post-hoc calibration corrects the *aggregate* relationship between stated confidence and empirical accuracy, but it cannot manufacture per-decision discrimination absent from the raw signal: the recalibration maps are monotonic, so the ranking of decisions is unchanged. On SWE-bench, corrected confidence still does not rank decisions by correctness because the underlying signal is uniformly high. We regard this distinction between aggregate calibration and per-decision discrimination as the paper’s central technical point for oversight design.

5.3 Self-consistency is a stronger signal on SWE-bench in a preliminary comparison

Figure 2 shows the self-consistency comparison on SWE-bench and GSM8K (Llama 3.3 70B, $n=110$, 3 samples each). On SWE-bench, verbal confidence yields lift = 1.47 and AUC-RC = 0.669; self-consistency yields lift = 2.06 and AUC-RC = 0.623. In this preliminary comparison, vote-fraction agreement is the substantially stronger error signal for the open-ended task. On GSM8K, verbal confidence is the stronger signal (lift = 2.86 vs. SC = 1.95), consistent with models being able to express genuine introspective uncertainty about arithmetic.

A plausible mechanism is that sampling probes the model’s implicit probability distribution over answers, making consistency a richer signal than a verbalised number produced in a single forward pass (Manakul et al., 2023; Kuhn et al., 2023). For governance frameworks requiring an oversight trigger on agent tasks, verbal confidence was insufficient in our experiments; multi-sample agreement provided a meaningfully better signal on this task at a $3\times$ API cost per decision, which may be acceptable for high-stakes items. Given the sample size, this result motivates further investigation rather than an immediate policy prescription.

5.4 Miscalibration concentrates on easy tasks

Stratifying Llama 70B on SWE-bench by difficulty: on issues labelled “<15 min fix” the model is 30% accurate while reporting 81% mean confidence (gap +0.51). On 15-min-to-4-hour issues the gap is under 0.10. The model treats nearly all issues as hard, failing precisely where an overseer would least expect it. By repository, overconfidence concentrates on *sphinx-doc* (gap +0.31) and *sympy* (+0.23).

6. Implementing the Audit-Trail Requirement

We implement the Table 1 schema as an append-only, hash-chained, HMAC-signed log using only the Python standard library. The construction itself is standard, drawing on classic secure-logging designs (Bellare & Yee, 1997; Schneier & Kelsey, 1998; Crosby & Wallach, 2009); the contribution is the mapping to the 2026 regulatory clauses, not the cryptography. Each entry stores *prev_hash* (SHA-256 of the previous entry’s canonical JSON) and a *signature* (HMAC-SHA256 over the entry content). A *verify()* function re-walks the chain and localises the first failure.

Two insider-attack simulations are run on a 150-entry chain: flipping *oversight_flagged* from *true* to *false* (suppressing a review trigger), and rewriting a *confidence* value upward (faking certainty). Both are detected and localised to the exact tampered entry; the intact chain verifies cleanly.

This is a prototype. HMAC is symmetric; production requires asymmetric signatures (e.g., Ed25519) and managed key rotation. What it demonstrates (append-chaining, per-entry signing, and tamper localisation) matches what the NIST/NCCoE concept paper asks for at the mechanism level.

7. Discussion

Why confidence fails on the agent-style task. SWE-bench elicits $p(\text{fix} > 15 \text{ min})$, an external difficulty prediction against a human annotation, not a self-assessment of the model’s correctness. Even a perfectly introspective model might have no useful prior over fix duration. The QA tasks elicit $p(\text{own answer is correct})$, a genuinely introspective quantity. Within this evaluation, the models tested appear capable of meaningful self-assessment on closed-ended reasoning but not on open-ended difficulty estimation. Both failure modes are governance-relevant; they call for different interventions.

Implications for framework design. A requirement that agents “flag low-confidence outputs” without specifying task type is likely to be implemented with verbal confidence gates, which worked on the structured tasks and failed on the open-ended judgment task in our evaluation. The specific recommendation is: governance requirements for confidence-based oversight should (a) specify whether they intend introspective confidence or difficulty estimation, (b) require empirical validation that the chosen signal provides oversight lift > 1 on the target task distribution, and (c) consider multi-sample agreement as a candidate alternative for open-ended tasks, pending replication at larger scale.

Limitations. The study uses single-turn judgment proxies for multi-step agent behaviour; results may not transfer to full agentic rollouts. Only two model families were tested, and results should not be generalised beyond them; reasoning models with chain-of-thought may have different self-assessment properties (Kadavath et al., 2022). Findings may also be sensitive to the elicitation prompt format, which was ablated for only one model. Sample sizes of 110–250 per condition produce CIs spanning 0.08–0.12 ECE units. The self-consistency comparison is preliminary ($n=110$, one model, SWE-bench and GSM8K only). The audit log uses symmetric signing. Finally, the regulatory landscape is moving: under the EU Digital Omnibus amendments agreed in 2026, the application of the AI Act’s high-risk obligations (Annex III) has been deferred to December 2027, so the operative timeline for Article 14 compliance is later than originally scheduled.

8. Conclusion

Three 2026 governance frameworks require confidence-based human oversight of AI agents but do not test whether the mechanism works. We report such a test for two widely deployed models across four benchmarks. For the models and tasks evaluated, verbalized confidence provides a useful oversight signal on structured QA (lift up to 2.9) but not on software-engineering difficulty judgment, an open-ended agent-style task (lift ≈ 1.0). Post-hoc recalibration repairs aggregate calibration but not per-decision discrimination. A preliminary self-consistency comparison ($n=110$) raises SWE-bench lift to 2.06. Within this scope, the results suggest that governance frameworks mandating confidence-gating without specifying task type, or without requiring

empirical validation on the target distribution, risk failing silently on the open-ended judgments where deployed agents commonly operate.

Data and code availability. All elicitation data, analysis scripts, figures, and the audit-log implementation are available at <https://github.com/Rushhaabhhh/confidence-gating-evaluation>. Experiments use fixed seeds and temperature-0 elicitation; every number and figure in this paper can be regenerated from the committed data without API access (`verify_claims.sh`).

References

- Bansal, G., Wu, T., Zhou, J., Fok, R., Nushi, B., Kamar, E., Ribeiro, M. T., & Weld, D. S. (2021). Does the whole exceed its parts? The effect of AI explanations on complementary team performance. *CHI 2021*.
- Bellare, M. & Yee, B. (1997). Forward integrity for secure audit logs. Technical report, University of California, San Diego.
- Brier, G. W. (1950). Verification of forecasts expressed in terms of probability. *Monthly Weather Review*, 78(1):1–3.
- Chan, A. et al. (2024). Visibility into AI Agents. *FAccT 2024*, pp. 958–973.
- Cobbe, K. et al. (2021). Training verifiers to solve math word problems. *arXiv:2110.14168*.
- Crosby, S. A. & Wallach, D. S. (2009). Efficient data structures for tamper-evident logging. *USENIX Security 2009*.
- Cloud Security Alliance (2026). *NIST AI Agent Standards Initiative*. CSA Lab Note.
- Regulation (EU) 2024/1689 (EU AI Act), Articles 14 and 15, as amended.
- Galileo (2026). *How to Build Human-in-the-Loop Oversight for AI Agents*.
- Geifman, Y. & El-Yaniv, R. (2017). Selective classification for deep neural networks. *NeurIPS 2017*.
- Gneiting, T. & Raftery, A. E. (2007). Strictly proper scoring rules, prediction, and estimation. *JASA*, 102(477):359–378.
- Guo, C., Pleiss, G., Sun, Y., & Weinberger, K. Q. (2017). On calibration of modern neural networks. *ICML 2017*.
- Hendrycks, D. et al. (2021). Measuring massive multitask language understanding. *ICLR 2021*.
- Infocomm Media Development Authority, Singapore (2026). *Model AI Governance Framework for Agentic AI*, version 1.5, May 2026.
- Jimenez, C. E. et al. (2024). SWE-bench: Can language models resolve real GitHub issues? *ICLR 2024*.
- Kadavath, S. et al. (2022). Language models (mostly) know what they know. *arXiv:2207.05221*.
- Kaddour, J., Patel, S., Dovonon, G., Richter, L., Minervini, P., & Kusner, M. J. (2026). Agentic uncertainty reveals agentic overconfidence. *arXiv:2602.06948*.
- Kuhn, L., Gal, Y., & Farquhar, S. (2023). Semantic uncertainty: linguistic invariances for uncertainty estimation in natural language generation. *ICLR 2023*.
- Lin, S., Hilton, J., & Evans, O. (2022). TruthfulQA: measuring how models mimic human falsehoods. *ACL 2022*.
- Manakul, P., Liusie, A., & Gales, M. J. F. (2023).

- SelfCheckGPT: zero-resource black-box hallucination detection for generative large language models. *EMNLP 2023*.
- NIST/NCCoE (2026). *Accelerating the Adoption of Software and AI Agent Identity and Authorization*. Concept paper, initial public draft, February 2026.
- Ouyang, L. et al. (2022). Training language models to follow instructions with human feedback. *NeurIPS 2022*.
- Parasuraman, R. & Manzey, D. H. (2010). Complacency and bias in human use of automation: an attentional integration. *Human Factors*, 52(3):381–410.
- Platt, J. (1999). Probabilistic outputs for support vector machines and comparisons to regularized likelihood methods. *Advances in Large Margin Classifiers*.
- Schneier, B. & Kelsey, J. (1998). Cryptographic support for secure logs on untrusted machines. *USENIX Security 1998*.
- Tian, K. et al. (2023). Just ask for calibration: strategies for eliciting calibrated confidence scores from language models fine-tuned with human feedback. *EMNLP 2023*.
- Vovk, V., Gammerman, A., & Shafer, G. (2005). *Algorithmic Learning in a Random World*. Springer.
- Wang, X. et al. (2023). Self-consistency improves chain-of-thought reasoning in language models. *ICLR 2023*.
- Xiong, M. et al. (2024). Can LLMs express their uncertainty? An empirical evaluation of confidence elicitation in LLMs. *ICLR 2024*.
- Zadrozny, B. & Elkan, C. (2002). Transforming classifier scores into accurate multiclass probability estimates. *KDD 2002*.

Appendix: Figures

Note. Figures are placed in the appendix to allow the main text to run uninterrupted. All figures are referenced in the results section; readers are encouraged to view them alongside the corresponding subsection.




Figure 1: Fig. 1 — Oversight Lift by Model and Domain. Lift = error rate among the least-confident 10% ÷ overall base error rate. Lift > 1 (green region): low confidence reliably finds errors; lift $\approx$ 1 (red region): the gate adds no value. SWE-bench lift is near 1 for both models across every configuration tested. GSM8K, MMLU, and TruthfulQA lifts are 1.67–2.86. *This is the paper’s primary result: for the two models evaluated, confidence-gating works on structured QA but not on the SWE-bench difficulty-judgment task.*




Figure 2: Fig. 2 — Self-Consistency vs. Verbal Confidence as Oversight Signals (Llama 3.3 70B, $n=110$). *Left:* oversight lift. *Right:* AUC-RC. On SWE-bench, self-consistency (3 samples, $T=0.3$) raises lift from 1.47 to 2.06, a 40% improvement on the task type where verbal confidence fails. On GSM8K, verbal confidence is the stronger signal (lift 2.86 vs. SC 1.95). Neither method dominates across task types, suggesting that the choice of uncertainty estimator should be matched to the task.




Figure 3: Fig. 3 — Reliability Diagrams (2 models $\times$ 4 domains). Bar height is empirical accuracy in each confidence bin; the dashed diagonal is perfect calibration. Red bars: overconfident (stated confidence exceeds accuracy). Blue: underconfident. GSM8K panels show systematic overconfidence for both models, with bars far above the diagonal. On MMLU and TruthfulQA, calibration is closer to the diagonal but bins are sparse due to confidence collapse.




Figure 4: Fig. 4 — Confidence Distributions by Domain and Model. Histograms split by outcome (green: correct; red: incorrect). Every panel has $\sigma < 0.10$ (labelled), indicating confidence collapse: the distributions of correct and incorrect decisions overlap almost entirely. Gemini’s GSM8K column contains all 250 observations in a single bar at 1.0 ($\sigma=0.000$). Because the distributions are indistinguishable, no fixed threshold can reliably separate reliable from unreliable decisions.




Figure 5: Fig. 5 — Risk-Coverage Curves (one panel per domain). The x -axis is the fraction of decisions handled autonomously (sorted most-confident first); the y -axis is accuracy on that covered set. Dotted horizontal lines show each model’s baseline accuracy. Curves that track the baseline, as both models do on SWE-bench, indicate that confidence does not rank decisions by correctness. On GSM8K, MMLU, and TruthfulQA, curves are meaningfully above baseline at low coverage.




Figure 6: Fig. 6 — Aggregate Calibration Before and After Post-Hoc Correction. Raw ECE (red) vs. best five-fold out-of-sample post-hoc ECE (green). P=Platt scaling, I=Isotonic regression. Every condition achieves >50% relative ECE reduction. *Aggregate calibration is recoverable without retraining; per-decision discrimination is not.* The Gemini-GSM8K green bar reaches near-zero because isotonic regression maps the single confidence value (1.0 for all items) to the mean accuracy—ECE = 0 by construction, not by calibration.